

In [1]:

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt,seaborn as sns
```


In [2]:

```
train_df=pd.read_csv(r"C:\Users\smb06\Downloads\Mobile_Price_Classification_test.csv")
train_df
```

Out[2]:

	id	battery_power	blue	clock_speed	dual_sim	fc	four_g	int_memory	m_dep	mobile_wt	...
0	1	1043	1	1.8	1	14	0	5	0.1	193	...
1	2	841	1	0.5	1	4	1	61	0.8	191	...
2	3	1807	1	2.8	0	1	0	27	0.9	186	...
3	4	1546	0	0.5	1	18	1	25	0.5	96	...
4	5	1434	0	1.4	0	11	1	49	0.5	108	...
...
995	996	1700	1	1.9	0	0	1	54	0.5	170	...
996	997	609	0	1.8	1	0	0	13	0.9	186	...
997	998	1185	0	1.4	0	1	1	8	0.5	80	...
998	999	1533	1	0.5	1	0	0	50	0.4	171	...
999	1000	1270	1	0.5	0	4	1	35	0.1	140	...

1000 rows × 21 columns



In [3]:

```
test_df=pd.read_csv(r"C:\Users\smb06\Downloads\Mobile_Price_Classification_train.csv")
test_df
```

Out[3]:

	battery_power	blue	clock_speed	dual_sim	fc	four_g	int_memory	m_dep	mobile_wt	n_cores
0	842	0	2.2	0	1	0	7	0.6	188	2
1	1021	1	0.5	1	0	1	53	0.7	136	3
2	563	1	0.5	1	2	1	41	0.9	145	5
3	615	1	2.5	0	0	0	10	0.8	131	6
4	1821	1	1.2	0	13	1	44	0.6	141	2
...
1995	794	1	0.5	1	0	1	2	0.8	106	6
1996	1965	1	2.6	1	0	0	39	0.2	187	4
1997	1911	0	0.9	1	1	1	36	0.7	108	8
1998	1512	0	0.9	0	4	1	46	0.1	145	5
1999	510	1	2.0	1	5	1	45	0.9	168	6

2000 rows × 21 columns

In [4]:

```
train_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 21 columns):
#   Column          Non-Null Count  Dtype
---  -
0   id              1000 non-null   int64
1   battery_power   1000 non-null   int64
2   blue            1000 non-null   int64
3   clock_speed     1000 non-null   float64
4   dual_sim        1000 non-null   int64
5   fc              1000 non-null   int64
6   four_g          1000 non-null   int64
7   int_memory      1000 non-null   int64
8   m_dep           1000 non-null   float64
9   mobile_wt       1000 non-null   int64
10  n_cores         1000 non-null   int64
11  pc              1000 non-null   int64
12  px_height       1000 non-null   int64
13  px_width        1000 non-null   int64
14  ram             1000 non-null   int64
15  sc_h            1000 non-null   int64
16  sc_w            1000 non-null   int64
17  talk_time       1000 non-null   int64
18  three_g         1000 non-null   int64
19  touch_screen    1000 non-null   int64
20  wifi            1000 non-null   int64
dtypes: float64(2), int64(19)
memory usage: 164.2 KB
```

In [5]:

```
test_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2000 entries, 0 to 1999
Data columns (total 21 columns):
 #   Column              Non-Null Count  Dtype  
---  -
 0   battery_power       2000 non-null   int64  
 1   blue                 2000 non-null   int64  
 2   clock_speed         2000 non-null   float64 
 3   dual_sim            2000 non-null   int64  
 4   fc                   2000 non-null   int64  
 5   four_g              2000 non-null   int64  
 6   int_memory          2000 non-null   int64  
 7   m_dep               2000 non-null   float64 
 8   mobile_wt           2000 non-null   int64  
 9   n_cores             2000 non-null   int64  
10   pc                   2000 non-null   int64  
11   px_height           2000 non-null   int64  
12   px_width            2000 non-null   int64  
13   ram                 2000 non-null   int64  
14   sc_h                2000 non-null   int64  
15   sc_w                2000 non-null   int64  
16   talk_time           2000 non-null   int64  
17   three_g             2000 non-null   int64  
18   touch_screen        2000 non-null   int64  
19   wifi                2000 non-null   int64  
20   price_range         2000 non-null   int64  
dtypes: float64(2), int64(19)
memory usage: 328.3 KB
```

In [6]:

```
x=train_df.drop('wifi',axis=1)
y=train_df['wifi']
```

In [7]:

```
x=test_df.drop('wifi',axis=1)
y=test_df['wifi']
```

In [8]:

```
train_df['dual_sim'].value_counts()
```

Out[8]:

```
dual_sim
1      517
0      483
Name: count, dtype: int64
```

In [9]:

```
test_df['blue'].value_counts()
```

Out[9]:

```
blue
0    1010
1     990
Name: count, dtype: int64
```

In [10]:

```
T={"Home Owner":{"Yes":1,"No":0}}
train_df=train_df.replace(T)
print(train_df)
```

	id	battery_power	blue	clock_speed	dual_sim	fc	four_g	int_memory	
0	1	1043	1	1.8	1	14	0	5	\
1	2	841	1	0.5	1	4	1	61	
2	3	1807	1	2.8	0	1	0	27	
3	4	1546	0	0.5	1	18	1	25	
4	5	1434	0	1.4	0	11	1	49	
..	
995	996	1700	1	1.9	0	0	1	54	
996	997	609	0	1.8	1	0	0	13	
997	998	1185	0	1.4	0	1	1	8	
998	999	1533	1	0.5	1	0	0	50	
999	1000	1270	1	0.5	0	4	1	35	

	m_dep	mobile_wt	...	pc	px_height	px_width	ram	sc_h	sc_w	
0	0.1	193	...	16	226	1412	3476	12	7	\
1	0.8	191	...	12	746	857	3895	6	0	
2	0.9	186	...	4	1270	1366	2396	17	10	
3	0.5	96	...	20	295	1752	3893	10	0	
4	0.5	108	...	18	749	810	1773	15	8	
..	
995	0.5	170	...	17	644	913	2121	14	8	
996	0.9	186	...	2	1152	1632	1933	8	1	
997	0.5	80	...	12	477	825	1223	5	0	
998	0.4	171	...	12	38	832	2509	15	11	
999	0.1	140	...	19	457	608	2828	9	2	

	talk_time	three_g	touch_screen	wifi
0	2	0	1	0
1	7	1	0	0
2	10	0	1	1
3	7	1	1	0
4	7	1	0	1
..
995	15	1	1	0
996	19	0	1	1
997	14	1	0	0
998	6	0	1	0
999	3	1	0	1

[1000 rows x 21 columns]

In [11]:

```
T={"Home Owner":{"Yes":1,"No":0}}
train_df=train_df.replace(T)
print(train_df)
```

	id	battery_power	blue	clock_speed	dual_sim	fc	four_g	int_memory	
0	1	1043	1	1.8	1	14	0	5	\
1	2	841	1	0.5	1	4	1	61	
2	3	1807	1	2.8	0	1	0	27	
3	4	1546	0	0.5	1	18	1	25	
4	5	1434	0	1.4	0	11	1	49	
..	
995	996	1700	1	1.9	0	0	1	54	
996	997	609	0	1.8	1	0	0	13	
997	998	1185	0	1.4	0	1	1	8	
998	999	1533	1	0.5	1	0	0	50	
999	1000	1270	1	0.5	0	4	1	35	

	m_dep	mobile_wt	...	pc	px_height	px_width	ram	sc_h	sc_w	
0	0.1	193	...	16	226	1412	3476	12	7	\
1	0.8	191	...	12	746	857	3895	6	0	
2	0.9	186	...	4	1270	1366	2396	17	10	
3	0.5	96	...	20	295	1752	3893	10	0	
4	0.5	108	...	18	749	810	1773	15	8	
..	
995	0.5	170	...	17	644	913	2121	14	8	
996	0.9	186	...	2	1152	1632	1933	8	1	
997	0.5	80	...	12	477	825	1223	5	0	
998	0.4	171	...	12	38	832	2509	15	11	
999	0.1	140	...	19	457	608	2828	9	2	

	talk_time	three_g	touch_screen	wifi
0	2	0	1	0
1	7	1	0	0
2	10	0	1	1
3	7	1	1	0
4	7	1	0	1
..
995	15	1	1	0
996	19	0	1	1
997	14	1	0	0
998	6	0	1	0
999	3	1	0	1

[1000 rows x 21 columns]

In [12]:

```
T={"Home Owner":{"Yes":1,"No":0}}
test_df=test_df.replace(T)
print(test_df)
```

	battery_power	blue	clock_speed	dual_sim	fc	four_g	int_memory	
0	842	0	2.2	0	1	0	7	\
1	1021	1	0.5	1	0	1	53	
2	563	1	0.5	1	2	1	41	
3	615	1	2.5	0	0	0	10	
4	1821	1	1.2	0	13	1	44	
...	
1995	794	1	0.5	1	0	1	2	
1996	1965	1	2.6	1	0	0	39	
1997	1911	0	0.9	1	1	1	36	
1998	1512	0	0.9	0	4	1	46	
1999	510	1	2.0	1	5	1	45	

	m_dep	mobile_wt	n_cores	...	px_height	px_width	ram	sc_h	sc_w	
0	0.6	188	2	...	20	756	2549	9	7	\
1	0.7	136	3	...	905	1988	2631	17	3	
2	0.9	145	5	...	1263	1716	2603	11	2	
3	0.8	131	6	...	1216	1786	2769	16	8	
4	0.6	141	2	...	1208	1212	1411	8	2	
...	
1995	0.8	106	6	...	1222	1890	668	13	4	
1996	0.2	187	4	...	915	1965	2032	11	10	
1997	0.7	108	8	...	868	1632	3057	9	1	
1998	0.1	145	5	...	336	670	869	18	10	
1999	0.9	168	6	...	483	754	3919	19	4	

	talk_time	three_g	touch_screen	wifi	price_range
0	19	0	0	1	1
1	7	1	1	0	2
2	9	1	1	0	2
3	11	1	0	0	2
4	15	1	1	0	1
...
1995	19	1	1	0	0
1996	16	1	1	1	2
1997	5	1	1	0	3
1998	19	1	1	1	0
1999	2	1	1	1	3

[2000 rows x 21 columns]

In [13]:

```
x=train_df.drop('wifi',axis=1)
y=train_df['wifi']
```

In [14]:

```
x=test_df.drop('wifi',axis=1)
y=test_df['wifi']
```

In [15]:

```
from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test = train_test_split(x,y,train_size=0.7,random_state=42)
x_train.shape,x_test.shape
```

Out[15]:

```
((1400, 20), (600, 20))
```

In [16]:

```
from sklearn.ensemble import RandomForestClassifier
rfc = RandomForestClassifier()
rfc.fit(x_train,y_train)
```

Out[16]:

```
▼ RandomForestClassifier
RandomForestClassifier()
```

In a Jupyter environment, please rerun this cell to show the HTML representation or

trust the notebook. On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

In [17]:

```
rf = RandomForestClassifier()
```

In [18]:

```
params = {'max_depth': [2,3,5,10,20],
'min_samples_leaf': [5,10,20,50,100,200],
'n_estimators': [10,25,30,50,100,200]}
```

In [19]:

```
from sklearn.model_selection import GridSearchCV
grid_search = GridSearchCV(estimator=rf,param_grid=params,cv = 2, scoring='accuracy')
grid_search.fit(x_train,y_train)
```

Out[19]:

```
GridSearchCV
└─ estimator: RandomForestClassifier
   └─ RandomForestClassifier
```

In a Jupyter environment, please rerun this cell to show the HTML representation or

In [20]:

```
grid_search.best_score_
```

Out[20]:

0.5307142857142857

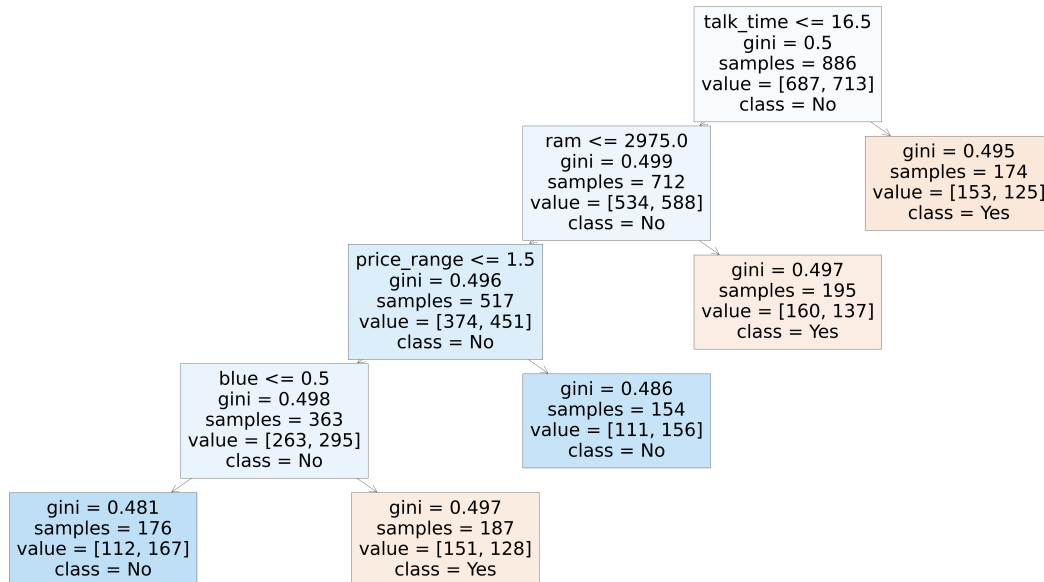
In [21]:

```
rf_best = grid_search.best_estimator_  
print(rf_best)
```

RandomForestClassifier(max_depth=5, min_samples_leaf=100, n_estimators=30)

In [23]:

```
from sklearn.tree import plot_tree  
plt.figure(figsize=(80,40))  
plot_tree(rf_best.estimators_[5], feature_names = x.columns, class_names=['Yes', "No"], filled=True)
```

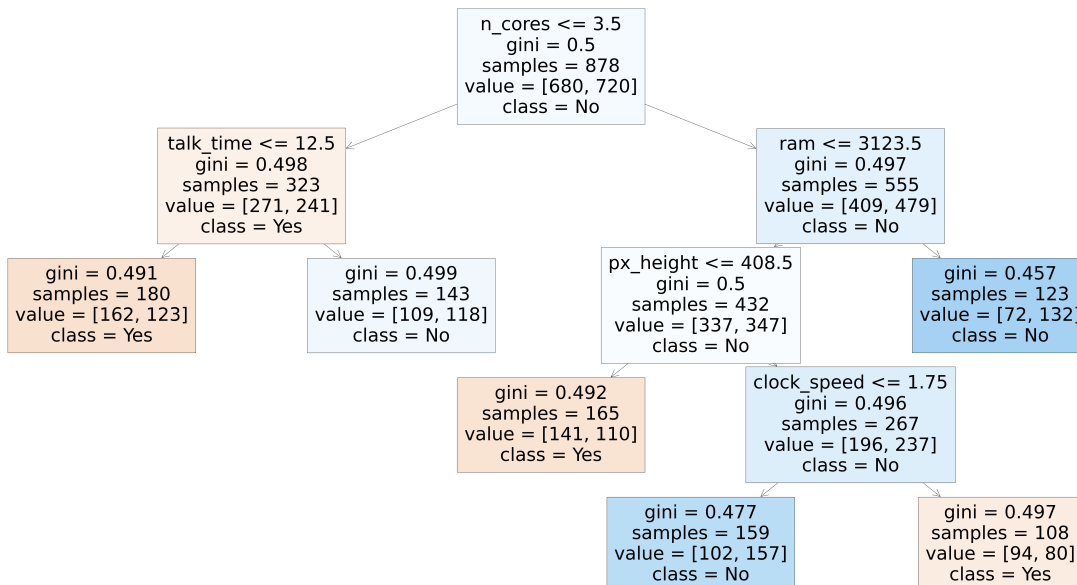


In [24]:

```
from sklearn.tree import plot_tree
plt.figure(figsize=(80,40))
plot_tree(rf_best.estimators_[7],feature_names=x.columns,class_names=["Yes","No"],filled=True)
```

Out[24]:

```
[Text(0.5, 0.9, 'n_cores <= 3.5\nngini = 0.5\nnsamples = 878\nnvalue = [680, 720]\nnclass = No'),
Text(0.25, 0.7, 'talk_time <= 12.5\nngini = 0.498\nnsamples = 323\nnvalue = [271, 241]\nnclass = Yes'),
Text(0.125, 0.5, 'gini = 0.491\nnsamples = 180\nnvalue = [162, 123]\nnclass = Yes'),
Text(0.375, 0.5, 'gini = 0.499\nnsamples = 143\nnvalue = [109, 118]\nnclass = No'),
Text(0.75, 0.7, 'ram <= 3123.5\nngini = 0.497\nnsamples = 555\nnvalue = [409, 479]\nnclass = No'),
Text(0.625, 0.5, 'px_height <= 408.5\nngini = 0.5\nnsamples = 432\nnvalue = [337, 347]\nnclass = No'),
Text(0.5, 0.3, 'gini = 0.492\nnsamples = 165\nnvalue = [141, 110]\nnclass = Yes'),
Text(0.75, 0.3, 'clock_speed <= 1.75\nngini = 0.496\nnsamples = 267\nnvalue = [196, 237]\nnclass = No'),
Text(0.625, 0.1, 'gini = 0.477\nnsamples = 159\nnvalue = [102, 157]\nnclass = No'),
Text(0.875, 0.1, 'gini = 0.497\nnsamples = 108\nnvalue = [94, 80]\nnclass = Yes'),
Text(0.875, 0.5, 'gini = 0.457\nnsamples = 123\nnvalue = [72, 132]\nnclass = No')]
```



In [25]:

```
rf_best.feature_importances_
```

Out[25]:

```
array([0.03334591, 0.01914167, 0.06293353, 0.00879701, 0.06890239,
       0.01881253, 0.06233112, 0.05192067, 0.07849322, 0.01252026,
       0.06905562, 0.14658583, 0.09846751, 0.05924079, 0.04058377,
       0.04552378, 0.08622158, 0.01038129, 0.00482762, 0.0219139 ])
```

In [26]:

```
imp_df = pd.DataFrame({"Vername": x_train.columns, "Imp": rf_best.feature_importances_})  
imp_df.sort_values(by="Imp", ascending=False)
```

Out[26]:

	Vername	Imp
11	px_height	0.146586
12	px_width	0.098468
16	talk_time	0.086222
8	mobile_wt	0.078493
10	pc	0.069056
4	fc	0.068902
2	clock_speed	0.062934
6	int_memory	0.062331
13	ram	0.059241
7	m_dep	0.051921
15	sc_w	0.045524
14	sc_h	0.040584
0	battery_power	0.033346
19	price_range	0.021914
1	blue	0.019142
5	four_g	0.018813
9	n_cores	0.012520
17	three_g	0.010381
3	dual_sim	0.008797
18	touch_screen	0.004828

In []: